



A GNU Radio-based Full Duplex Radio System

***Adam Parower
The Aerospace Corporation***

September 13, 2017

Agenda

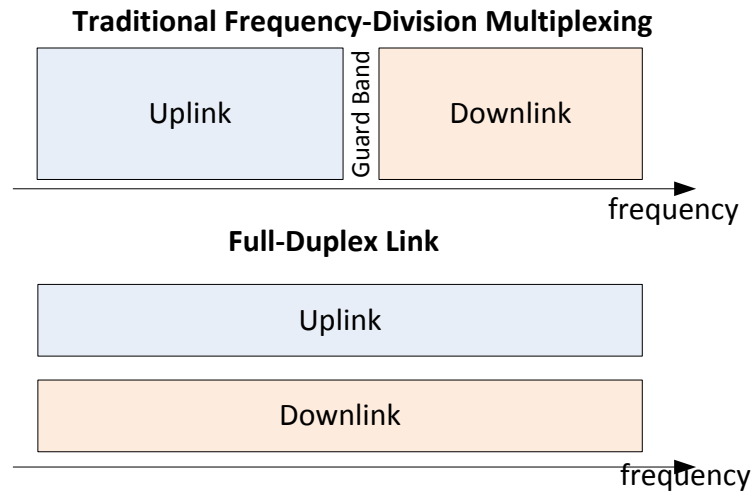
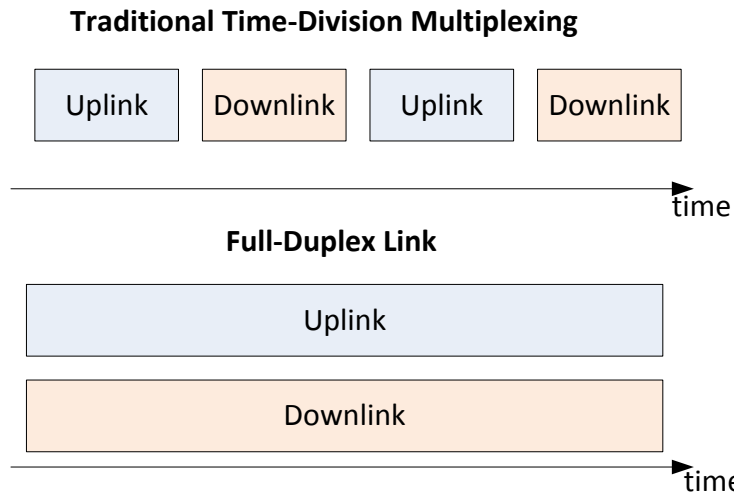


- Theory: Full-Duplex
 - *What is Full Duplex?*
 - *The Problem*
 - *The Solution*
- Theory: Digital Cancellation
 - *Black Box Model*
 - *Modeling Nonlinearity*
 - *Least-squares Problem*
 - *Algorithms for Solving*
- Custom GNU Radio Blocks
 - *Implementation*
 - *Cancellation Performance*
 - *Throughput Results*
- Using USRP Products
 - *Timing Synchronization*
 - *Built-in vs. External Mixers*
- Conclusions

What is Full-Duplex Communications?

... it's not the 'usual definition'


- Transmit and receive on the same antenna on the same frequency at the same time
 - Potentially double (or more) throughput in the same spectrum
 - Considered impossible for typical communications links
 - Multiple classical textbooks explicitly state it can't be done
- Can replace both time- and frequency-multiplexed links



The Problem



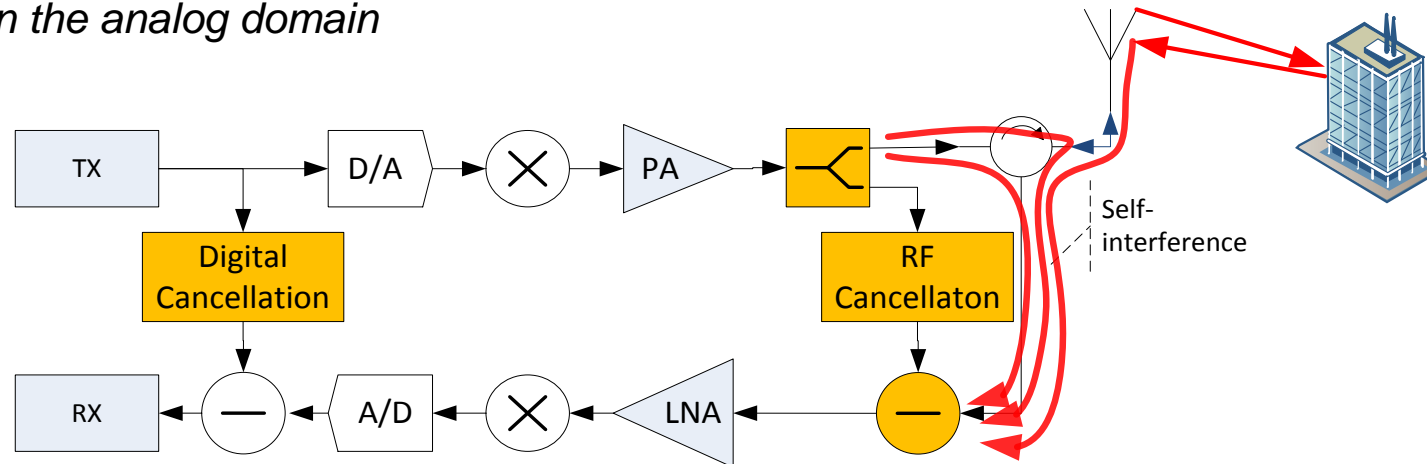
- The problem is self-interference
 - Transmit power swamps the receive power, making it very difficult to detect the desired signal
 - Power difference $10\log_{10}(P_{Tx}/P_{Rx})$ depends on the distance between Tx and Rx



The Solution

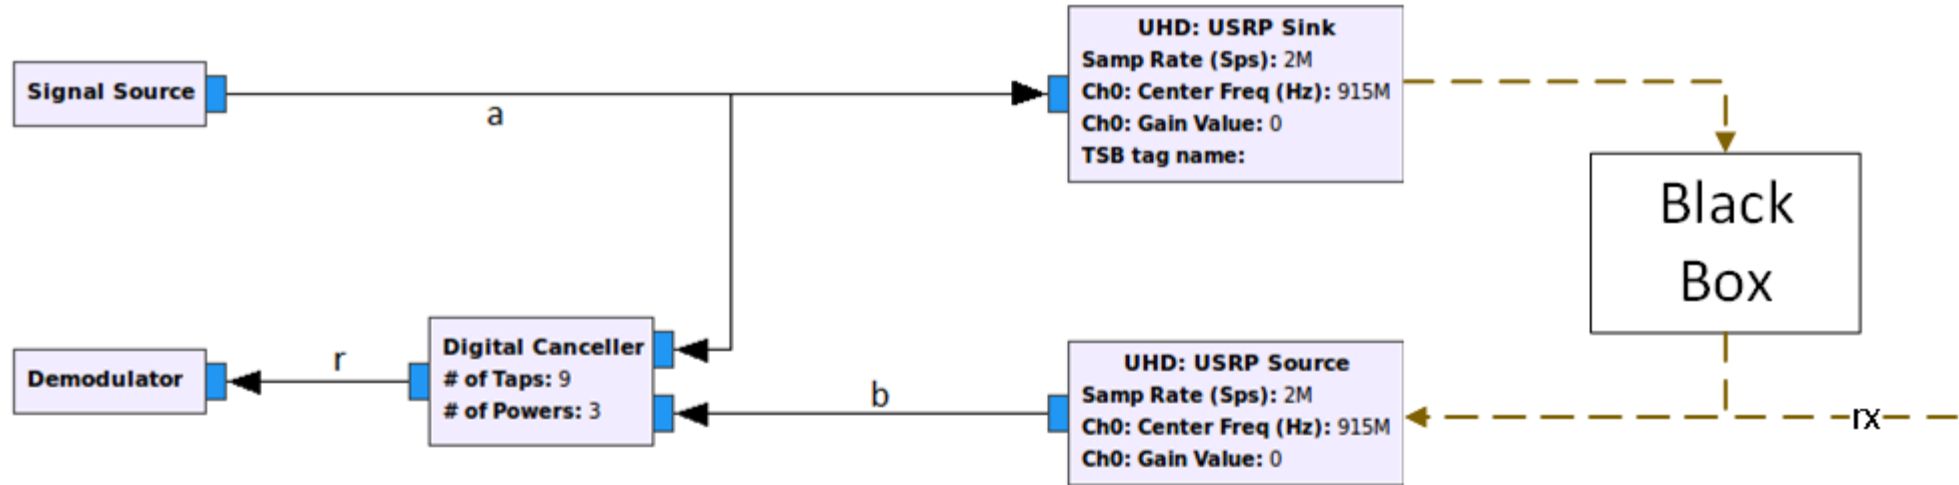
Very active area of research in the past several years

- The transmitted signal is known at the receiver
 - *To some degree, need to account for delay and distortion(s)*
- Subtract the transmitted signal from the received signal to remove self-interference
- The devil is in the implementation
 - *Transmitted signal power may be > 100 dB above received signal*
 - *Need very high linearity, very accurate matching and model of distortions*
- Two major approaches:
 1. *Cancel at baseband in the digital domain*
 2. *Cancel at RF in the analog domain*



All reported full-duplex systems use both analog and digital cancellation

Digital Cancellation



- Find a model for the black box to minimize r .
 - I.e. find \mathcal{F} to minimize $|r| = |b - \mathcal{F}(a)|$
- The model must account for:
 - Gain
 - Phase shift
 - Time delay
 - Nonlinear distortion
 - Multipath

Modeling the Black Box

Linear Model

- We model the black box as a non-causal FIR filter.

- Let:

a_i = transmitted signal, sample i

b_i = received signal, sample i

x_k = FIR filter coefficient k

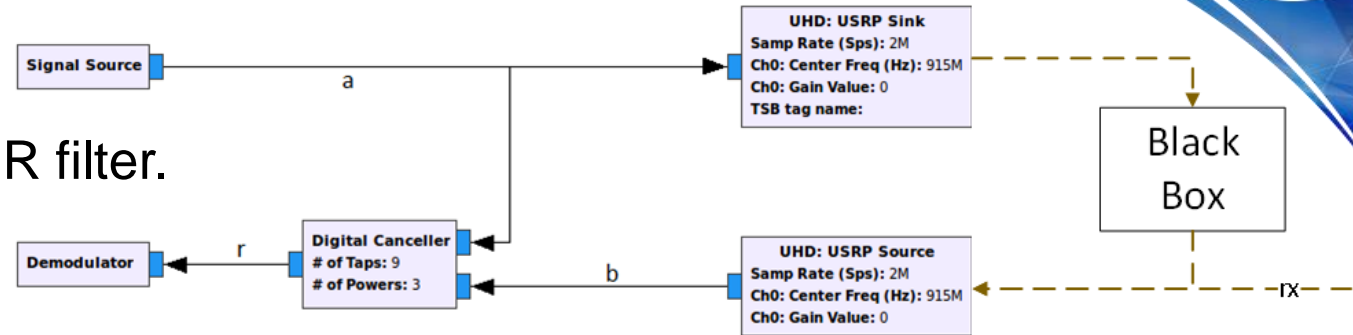
- Our goal is to find the filter coefficients that produce outputs (r_i) most similar to the actual received signal.

- Example for a 5-tap filter, for the 2nd received sample:

$$b_2 = x_0 a_0 + x_1 a_1 + x_2 a_2 + x_3 a_3 + x_4 a_4$$

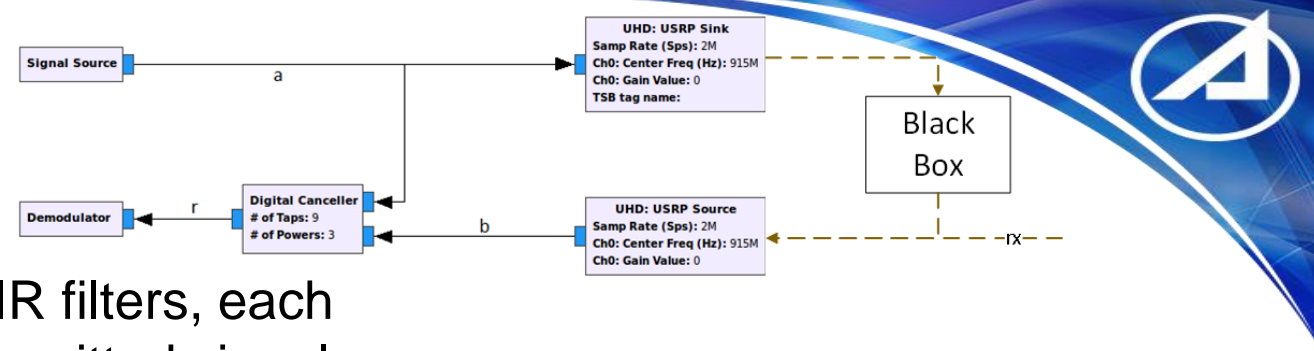
- Or, in vector form:

$$[a_0 \ a_1 \ a_2 \ a_3 \ a_4] \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = b_2$$



Modeling the Black Box

Nonlinear Distortion



- We model the black box as a **summation** of FIR filters, each operating on a different (odd) power of the transmitted signal.

- Let:

a_i = transmitted signal, sample i

b_i = received signal, sample i

x_{jk} = FIR filter coefficient k for power j

- Example for 3 powers (1st, 3rd, and 5th), each with 5 taps:

$$\begin{bmatrix} a_0 & a_1 & a_2 & a_3 & a_4 & a_0^3 & a_1^3 & a_2^3 & a_3^3 & a_4^3 & a_0^5 & a_1^5 & a_2^5 & a_3^5 & a_4^5 \end{bmatrix} \begin{bmatrix} x_{10} \\ x_{11} \\ x_{12} \\ x_{13} \\ x_{14} \\ x_{30} \\ x_{31} \\ x_{32} \\ x_{33} \\ x_{34} \\ x_{50} \\ x_{51} \\ x_{52} \\ x_{53} \\ x_{54} \end{bmatrix} = b_2$$



Modeling the Black Box

Matrix Representation

- If we wish to work with multiple samples of the received signal at a time, we can express the model in matrix form as follows:

$$\begin{bmatrix}
 a_0 & a_1 & a_2 & a_3 & a_4 & a_0^3 & a_1^3 & a_2^3 & a_3^3 & a_4^3 & a_0^5 & a_1^5 & a_2^5 & a_3^5 & a_4^5 \\
 a_1 & a_2 & a_3 & a_4 & a_5 & a_1^3 & a_2^3 & a_3^3 & a_4^3 & a_5^3 & a_1^5 & a_2^5 & a_3^5 & a_4^5 & a_5^5 \\
 a_2 & a_3 & a_4 & a_5 & a_6 & a_2^3 & a_3^3 & a_4^3 & a_5^3 & a_6^3 & a_2^5 & a_3^5 & a_4^5 & a_5^5 & a_6^5 \\
 a_3 & a_4 & a_5 & a_6 & a_7 & a_3^3 & a_4^3 & a_5^3 & a_6^3 & a_7^3 & a_3^5 & a_4^5 & a_5^5 & a_6^5 & a_7^5 \\
 a_4 & a_5 & a_6 & a_7 & a_8 & a_4^3 & a_5^3 & a_6^3 & a_7^3 & a_8^3 & a_4^5 & a_5^5 & a_6^5 & a_7^5 & a_8^5 \\
 a_5 & a_6 & a_7 & a_8 & a_9 & a_5^3 & a_6^3 & a_7^3 & a_8^3 & a_9^3 & a_5^5 & a_6^5 & a_7^5 & a_8^5 & a_9^5 \\
 & & & & & \vdots & & & & & & & & &
 \end{bmatrix}
 \begin{bmatrix}
 x_{10} \\
 x_{11} \\
 x_{12} \\
 x_{13} \\
 x_{14} \\
 x_{30} \\
 x_{31} \\
 x_{32} \\
 x_{33} \\
 x_{34} \\
 x_{50} \\
 x_{51} \\
 x_{52} \\
 x_{53} \\
 x_{54}
 \end{bmatrix}
 =
 \begin{bmatrix}
 b_2 \\
 b_3 \\
 b_4 \\
 b_5 \\
 b_6 \\
 b_7 \\
 \vdots
 \end{bmatrix}$$

$$A \vec{x} = \vec{b}$$

We are looking for the value of \vec{x} that minimizes $\|A\vec{x} - \vec{b}\|$

Algorithms

For Solving $A\vec{x} = \vec{b}$

Block-based algorithms operate on the entire matrix A , effectively computing $\vec{x} = A^+ \vec{b}$.

- Singular value decomposition (SVD)
- QR decomposition
- Cholesky decomposition
- Conjugate gradient method

Adaptive algorithms operate on one row of A at a time, adjusting the value of \vec{x} each iteration.

- Least mean squares (LMS)
- Recursive least squares (RLS)

Algorithms

For Solving $A\vec{x} = \vec{b}$

Block-based algorithms operate on the entire matrix A , effectively computing $\vec{x} = A^+ \vec{b}$.

- Singular value decomposition (SVD)
- QR decomposition
- Cholesky decomposition
- Conjugate gradient method

Adaptive algorithms operate on one row of A at a time, adjusting the value of \vec{x} each iteration.

- Least mean squares (LMS)
- Recursive least squares (RLS)

We implemented a GNU Radio block for each algorithm in green.

Note



- These GNU Radio blocks are not just for full duplex!
- They can be used to solve any problem where it is desirable to cancel a known signal that may have undergone linear and/or nonlinear distortion.



Performance Optimization

- For high throughput performance, we leverage the Intel libraries:
 - **IPP**: *Intel Performance Primitives*:
 - Contains signal processing routines
 - Optimized using Streaming SIMD Extensions
 - **MKL**: *Math Kernel Library*
 - Contains optimized functions for vector and matrix math
 - API is compatible with BLAS and LAPACK functions
- We further increase throughput by dividing computationally-intensive work between multiple threads.



Implementation (Pseudocode)

SVD Block

```
int svd_canceller_cc_impl::general_work(
    int noutput_items, gr_vector_int &ninput_items,
    gr_vector_const_void_star &input_items, gr_vector_void_star &output_items)
{
    gr_complex* ref = input_signals[0];    // ref = transmitted signal
    gr_complex* b   = input_signals[1];    // b = received signal
    gr_complex* out = output_signals[0];

    construct_A_matrix(ref, A);
    cgelsd(A, b, x, ...);                  // solve  $Ax = b$  for  $x$  (using SVD)
    cgemv(A, x, b, out, ...);              // residual =  $b - Ax$ 

    consume_each(block_size);
    return block_size;
}
```




Implementation (Pseudocode)

QR Decomposition Block

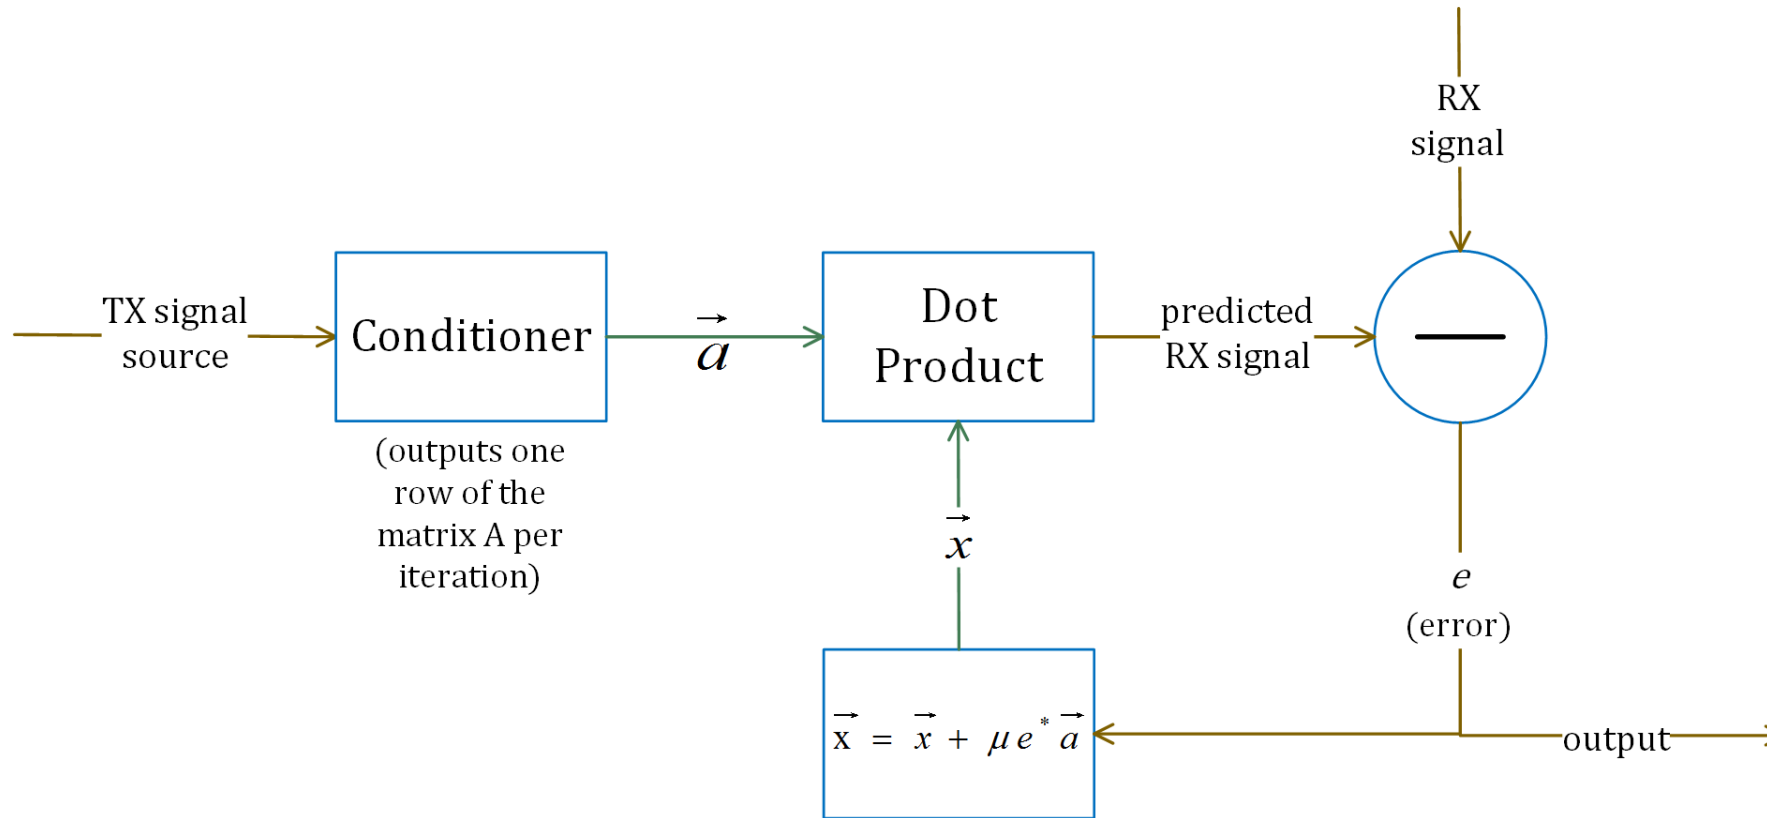
```
int qr_canceller_cc_impl::general_work(
    int noutput_items, gr_vector_int &ninput_items,
    gr_vector_const_void_star &input_items, gr_vector_void_star &output_items)
{
    gr_complex* ref = input_signals[0];    // ref = transmitted signal
    gr_complex* b   = input_signals[1];    // b = received signal
    gr_complex* out = output_signals[0];

    construct_A_matrix(ref, A);
    cgeqrf(A, temp, ...);                  // perform QR factorization on A
    cunqqr(temp, Q, ...);                  // compute Q explicitly
    cgemv(Q, b, Qb, ...);                  // compute  $Q^T b$ 
    cgemv(Q, Qb, b, out, ...);             // residual =  $b - Q Q^T b$ 

    consume_each(block_size);
    return block_size;
}
```

Implementation (Block Diagram)


LMS Block



- We optimize this block for throughput by:
 - Using multiple threads, each of which processes a chunk of the input
 - Calculating updates to the coefficients $\vec{\bar{x}}$ every N samples (instead of every sample)
 - Averaging the threads' values of $\vec{\bar{x}}$ when they synchronize

Results and Performance

Cancellation

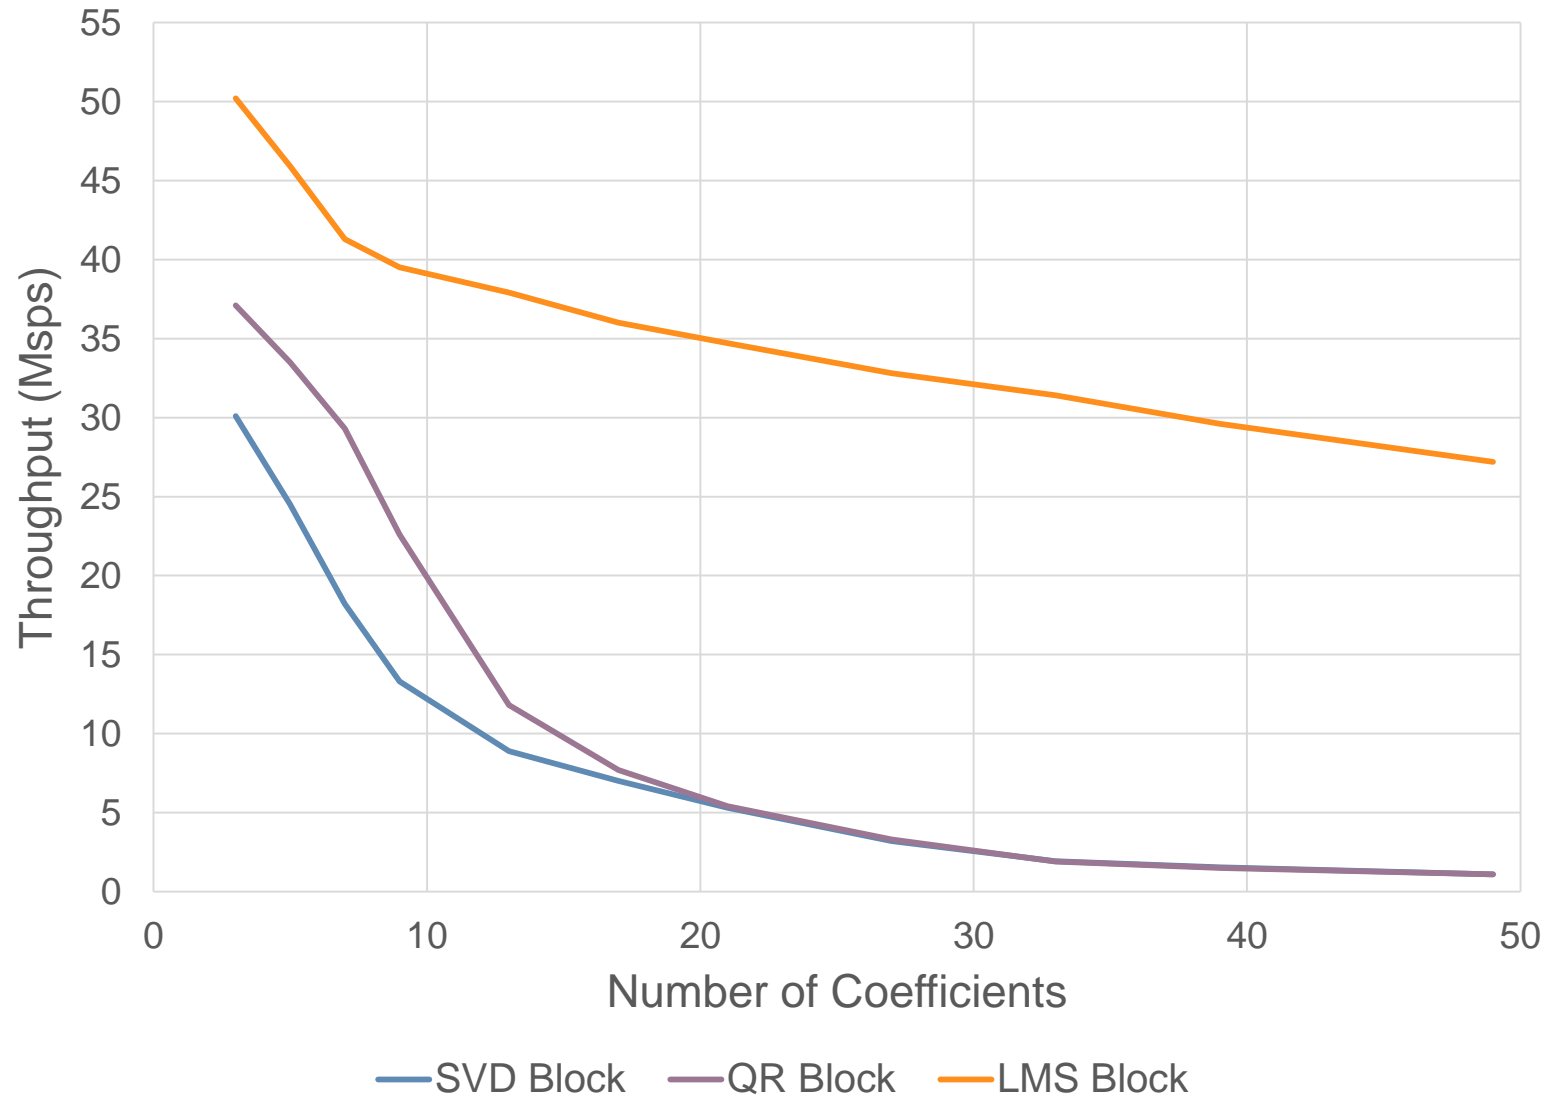


Test conditions:

- Digital simulation
- 13-tap FIR filter
- 1 power (linear only)
- SVD/QR block size: 512
- LMS step size: 0.01

Results and Performance

Throughput



Test conditions:

- Intel Xeon X5660 (2.80GHz)
- 8 threads

Using USRP Products

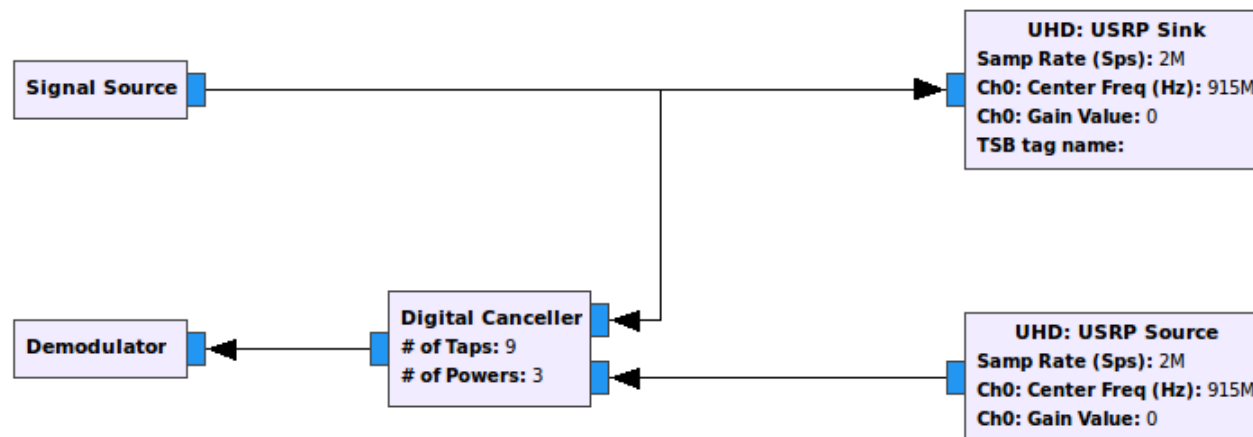
- If you decide to use USRPs in a full-duplex radio system, here are some tips ...



Source: https://www.ettus.com/content/images/prod_un210-kit_01_lg.jpg

Using USRP Products

Timing Synchronization



- For this flowgraph to work, the USRP sink and source blocks must start streaming at the same time.
- These blocks provide functions for timing synchronization:
 - `set_time_now()`
 - `set_time_next_pps()`
 - `set_start_time()`
- We can edit the GRC-generated Python code to call these functions ...
 - *But if we make a change to the flowgraph and regenerate, we have to do it again*
 - *There must be a better way ...*

Function Caller Block



- Wouldn't it be nice if there was a GRC block that allowed us to embed arbitrary function calls in the generated code, that execute before the flowgraph starts?

- We created one:

Function Caller ID: caller3 Function Belongs to: Block Block ID: usrp_source Function Name: set_time_now Function Args: 0	Function Caller ID: caller2 Function Belongs to: Block Block ID: usrp_sink Function Name: set_time_now Function Args: 0
Function Caller ID: caller1 Function Belongs to: Block Block ID: usrp_source Function Name: set_start_time Function Args: 1	Function Caller ID: caller0 Function Belongs to: Block Block ID: usrp_sink Function Name: set_start_time Function Args: 1

Generated Python code:

```
val = self.usrp_source.set_time_now(0)
...
val = self.usrp_sink.set_time_now(0)
...
val = self.usrp_source.set_start_time(1)
...
val = self.usrp_sink.set_start_time(1)
```

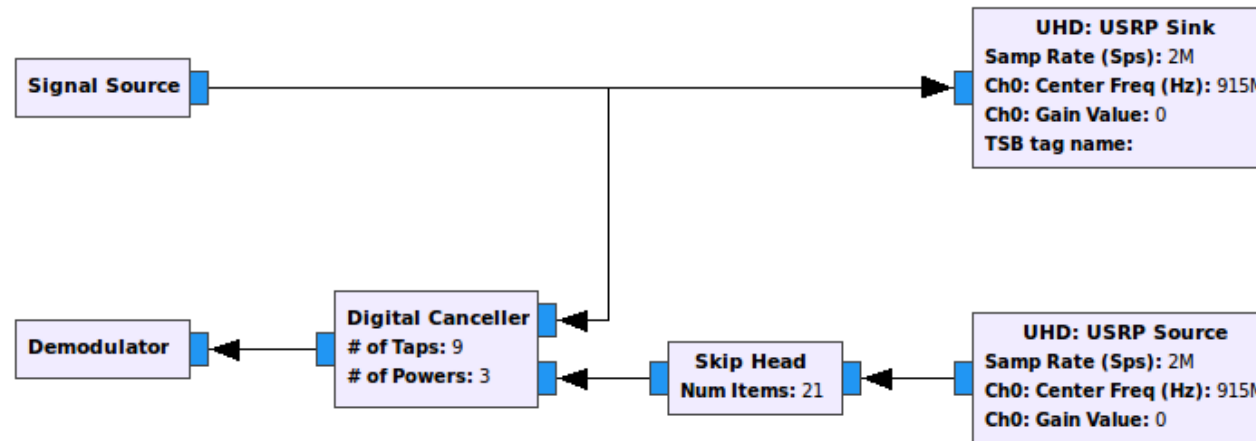
- Largely based on the built-in Function Probe block



Using USRP Products

Timing Synchronization




- Even after synchronizing the USRP source and sink, there is still a timing offset:
 - *Deterministic and repeatable*
 - *Appears to be sample-rate dependent*
 - *On the order of 10-50 samples*
- This can be remedied with a skip head block:



Using USRP Products

Built-in vs. External Mixers

- Most USRP daughterboards contain LO generators and mixers to convert between baseband/IF and RF.
- e.g. SBX, UBX
- Some daughterboards contain no LO or mixer but operate at baseband/IF.
- e.g. BasicTX, BasicRX



Source: <https://www.ettus.com/product/category/Daughterboards>

- These boards can be used with external mixers to produce RF.



Using USRP Products

Built-in vs. External Mixers

- For applications like full duplex where high SNR is essential, use external mixers for optimal performance:

Digital Cancellation (dB)		Transmitter Setup	
		BasicTX + External Up-converter	SBX
Receiver Setup	BasicRX + External Down-converter	54.4	44.6
	SBX	39.2	38.0

Test conditions:

- USRP X310
- Analog loopback



Summary and Conclusions

- It is feasible to implement a full-duplex radio system using GNU Radio.
- By using Intel libraries and multi-threading, we can support bandwidths in the tens of MHz.
- If parallelized, adaptive algorithms like LMS provide higher throughput with minimal cost to cancellation.
- Digital cancellation blocks can be applied to any scenario involving suppression of a known signal.

Questions?

